

Full Stack Developer Intern

AI-fluent — Internship at Gudz

About Gudz

Gudz builds modern supply chain software for new-age commerce. We come from inside the trenches — our team has spent years running logistics operations for some of India's fastest-moving retailers and quick-commerce players, including Zepto, BigBasket, DealShare and CityMall.

That ground-up exposure taught us where supply chains break, where data goes stale, and why most operators are stuck stitching together spreadsheets and disconnected dashboards. We're now building what we wished we had: a single source of truth product for supply chain operators — fast, opinionated, and built for the way modern commerce actually runs.

We're a small team that ships quickly, uses AI aggressively in our own workflows, and cares more about taste and outcomes than process.

The Role

We're hiring a Full Stack Developer Intern who is comfortable across the stack and excited about building with the latest generation of AI tools — not just calling an API, but designing real product features around LLMs, agents, and retrieval systems.

You'll work directly with the founders and core engineering team on features that go to production fast. Expect to ship in weeks, not quarters, and to own meaningful surfaces of the product.

What You'll Do

- Build full-stack features end-to-end — from database schema and API to the UI users actually click on.
- Design and ship AI-powered features: LLM-backed workflows, agents that take actions across our system, retrieval-augmented (RAG) experiences over operational data.
- Integrate with internal and external APIs — logistics partners, ERPs, marketplaces, payment and messaging platforms.
- Model data for supply chain entities (orders, SKUs, inventory, shipments) and write performant queries against them.
- Use AI dev tools (Claude Code, Cursor, v0, etc.) as a real productivity multiplier, not a gimmick.
- Participate in product decisions — we expect interns to push back, suggest better solutions, and have opinions on UX.
- Write code that other people can read, test what matters, and deploy without drama.

What We're Looking For

Must-haves

- Solid grasp of JavaScript/TypeScript and a modern frontend framework (React or Next.js preferred).
- Backend experience in Node.js, Python, or similar — building REST/GraphQL APIs, handling auth, deploying services.

- Working knowledge of relational databases (Postgres/MySQL) and at least basic SQL fluency. Familiarity with one NoSQL store (MongoDB, Redis) is a plus.
- Hands-on experience using LLM APIs (OpenAI, Anthropic Claude, Gemini, or open-source models) in a real project — even a side project counts.
- Comfort with Git, code reviews, and shipping to a deployed environment (Vercel, Render, Railway, AWS, GCP — anything).
- Strong product sense and a bias toward shipping. You should be able to look at a half-baked spec and ask the right questions before writing code.

Bonus Points

- Built something with agents, tool-calling, or function-calling beyond a basic chatbot.
- Experience with vector databases (Pinecone, Weaviate, pgvector) and RAG pipelines.
- Worked with workflow/orchestration frameworks (LangChain, LlamaIndex, Vercel AI SDK, Mastra) — or strong opinions on why you don't need them.
- Exposure to logistics, supply chain, e-commerce, or operations-heavy products.
- Comfort with infra basics: Docker, CI/CD, queues, background jobs.
- Open-source contributions, a public GitHub, or a personal portfolio of shipped projects.

AI Tooling Fluency

We expect every engineer at Gudz to use AI tools daily. You should already be using — or be eager to use — tools like Claude Code, Cursor, v0, Lovable, GitHub Copilot, or Replit Agent in your workflow. We don't care which one; we care that you have a real point of view on where they help and where they don't.

Who You Are

- A current student or recent graduate (B.Tech / BCA / MCA / equivalent), or a self-taught developer with the work to back it up.
- Curious by default — you read changelogs, try new tools the week they launch, and have strong instincts on what's hype vs. real.
- Comfortable in an early-stage environment where priorities shift, scope is fluid, and the answer to “how do we do X?” is often “figure it out.”
- Available to work a hybrid schedule from our office for 3–6 months.

What You'll Get

- Stipend of ₹18,000–₹25,000/month, calibrated to your skillset and assessed during the interview process.
- Direct mentorship from the founding team and senior engineers — no buffering through layers of management.
- Ownership of real product surfaces that go to paying customers, not throwaway intern projects.
- Access to paid AI tools and the latest models — we don't skimp on the stack.
- Strong consideration for a full-time offer at the end of the internship for high performers.

How to Apply

Send us your resume along with links to anything you've built — GitHub, deployed projects, side experiments, agent demos, or anything else that shows how you think and ship. A short note on the most interesting AI-powered thing you've built (or want to build) goes a long way. Applications can be submitted via your campus recruitment cell, or directly at: [insert application email / link].